



《图卷积神经网络》

主讲： 彭老师

<https://www.julyedu.com/>





图卷积神经网络

Graph Convolution Networks



GNN的基准化

Benchmarking Graph Neural Networks



/01 Graph Convolution Networks

图卷积神经网络



Dataset

- KarateClub: 数据为无向图，来源于论文[An Information Flow Model for Conflict and Fission in Small Groups](#);
- TUDataset: 包括58个基础的分类数据集集合，数据都为无向图，如“IMDB-BINARY”，“PROTEINS”等，来源于[TU Dortmund University](#)
- Planetoid: 引用网络数据集，包括“Cora”, “CiteSeer” and “PubMed”，数据都为无向图，来源于论文[Revisiting Semi-Supervised Learning with Graph Embeddings](#)。节点代表文档，边代表引用关系。
- CoraFull: 完整的“Cora” 引用网络数据集，数据为无向图，来源于论文[Deep Gaussian Embedding of Graphs: Unsupervised Inductive Learning via Ranking](#)。节点代表文档，边代表引用关系。

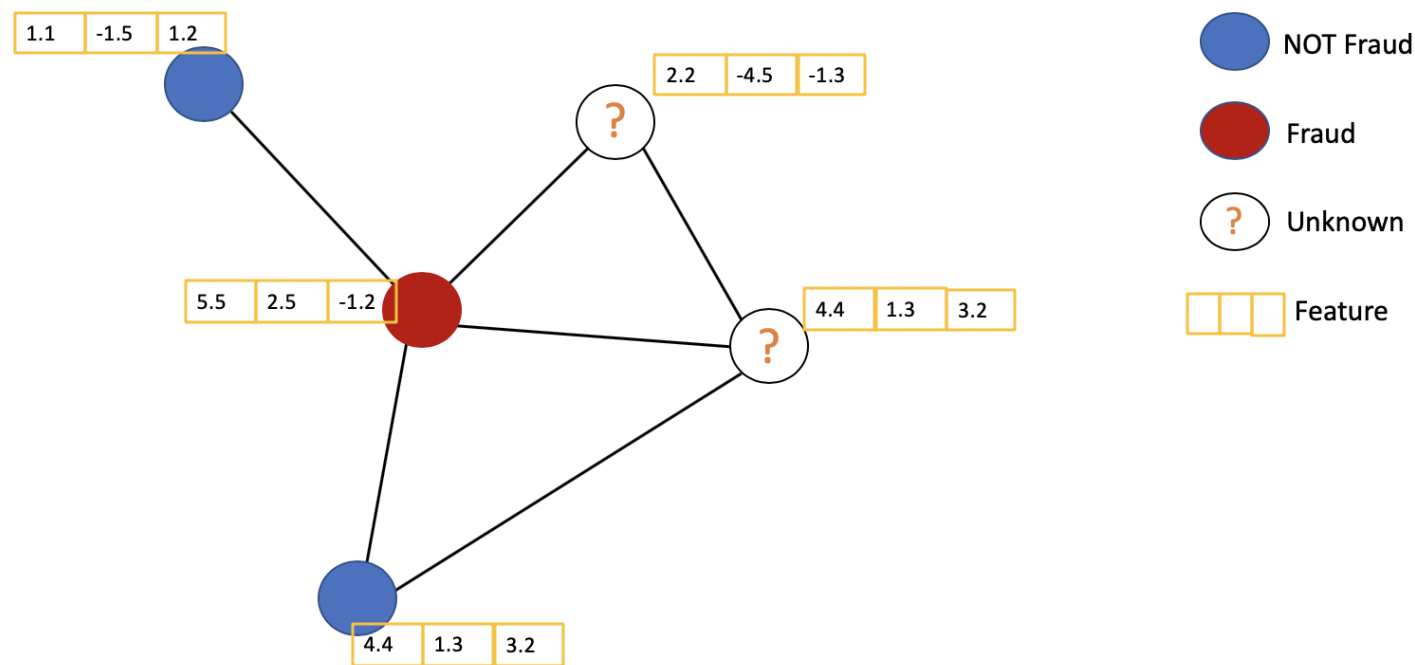
Dataset

- Coauthor: 共同作者网络数据集, 包括“CS”和“Physics”, 数据都为无向图, 来源于论文[Pitfalls of Graph Neural Network Evaluation](#)。节点代表作者, 若是共同作者则被边相连。学习任务是将作者映射到各自的研究领域中。
- Amazon: 亚马逊网络数据集, 包括“Computers”和“Photo”, 数据都为无向图, 来源于论文[Pitfalls of Graph Neural Network Evaluation](#)。节点代表货物, 边代表两种货物经常被同时购买。学习任务是将货物映射到各自的种类里。
- PPI: 蛋白质-蛋白质反应网络, 数据为无向图, 来源于论文[Predicting multicellular function through multi-layer tissue networks](#)
- Entities: 关系实体网络, 包括“AIFB”, “MUTAG”, “BGS”和“AM”, 数据都为无向图, 来源于论文[Modeling Relational Data with Graph Convolutional Networks](#)
- BitcoinOTC: 数据为有向图, 包括138个“who-trusts-whom”网络, 来源于论文[EvolveGCN: Evolving Graph Convolutional Networks for Dynamic Graphs](#), 数据链接为[Bitcoin OTC trust weighted signed network](#)

Graph Convolution Networks


Graph Convolution Networks (**GCN**) is a type of convolutional neural network that can work directly on graphs and take advantage of their structural information.

GCN solves the problem of **node classification** (such as documents) in a graph (such as a citation network), where labels are only available for a small subset of nodes (semi-supervised learning).



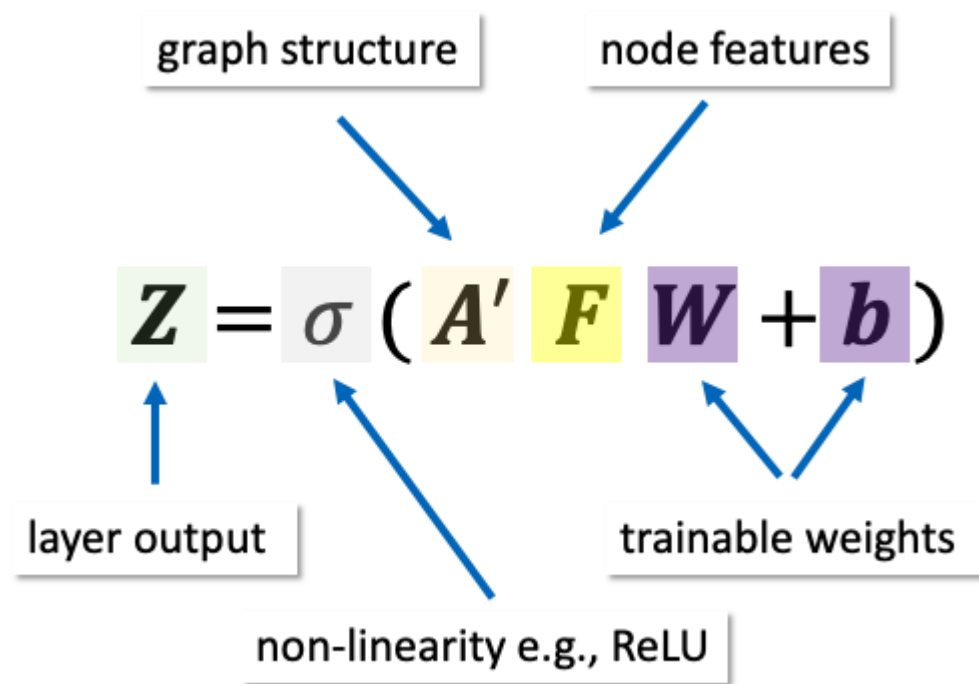
Comparison between CNN and GNN

The layer of a Graph Convolutional Neural Network (GCN) :





Comparison between CNN and GNN

The layer of a Graph Convolutional Neural Network (GCN) :



Single CNN layer with 3x3 filter:




$$H^{(l+1)} = \sigma\left(\tilde{D}^{-\frac{1}{2}} \tilde{A} \tilde{D}^{-\frac{1}{2}} H^{(l)} W^{(l)}\right)$$

Comparison between CNN and GNN


Receptive Field:

Left: CNN

Right: GNN



a)




b)

Comparison between CNN and GNN

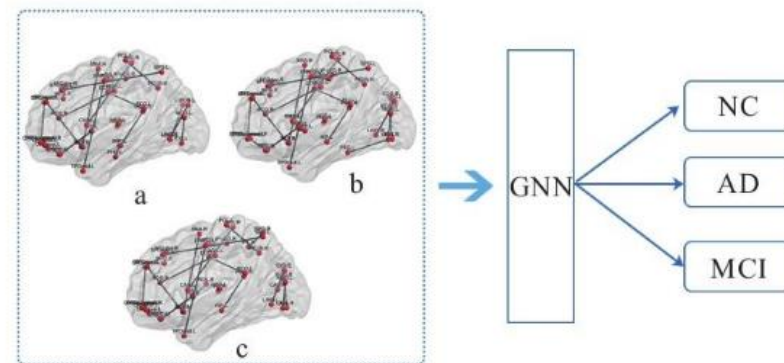
Tasks:

Left: CNN


Right: GNN




a) 图像分类



c) 图分类 NC: 正常 AD: 阿尔茨海默症患者 MCI: 轻度认知障碍患者

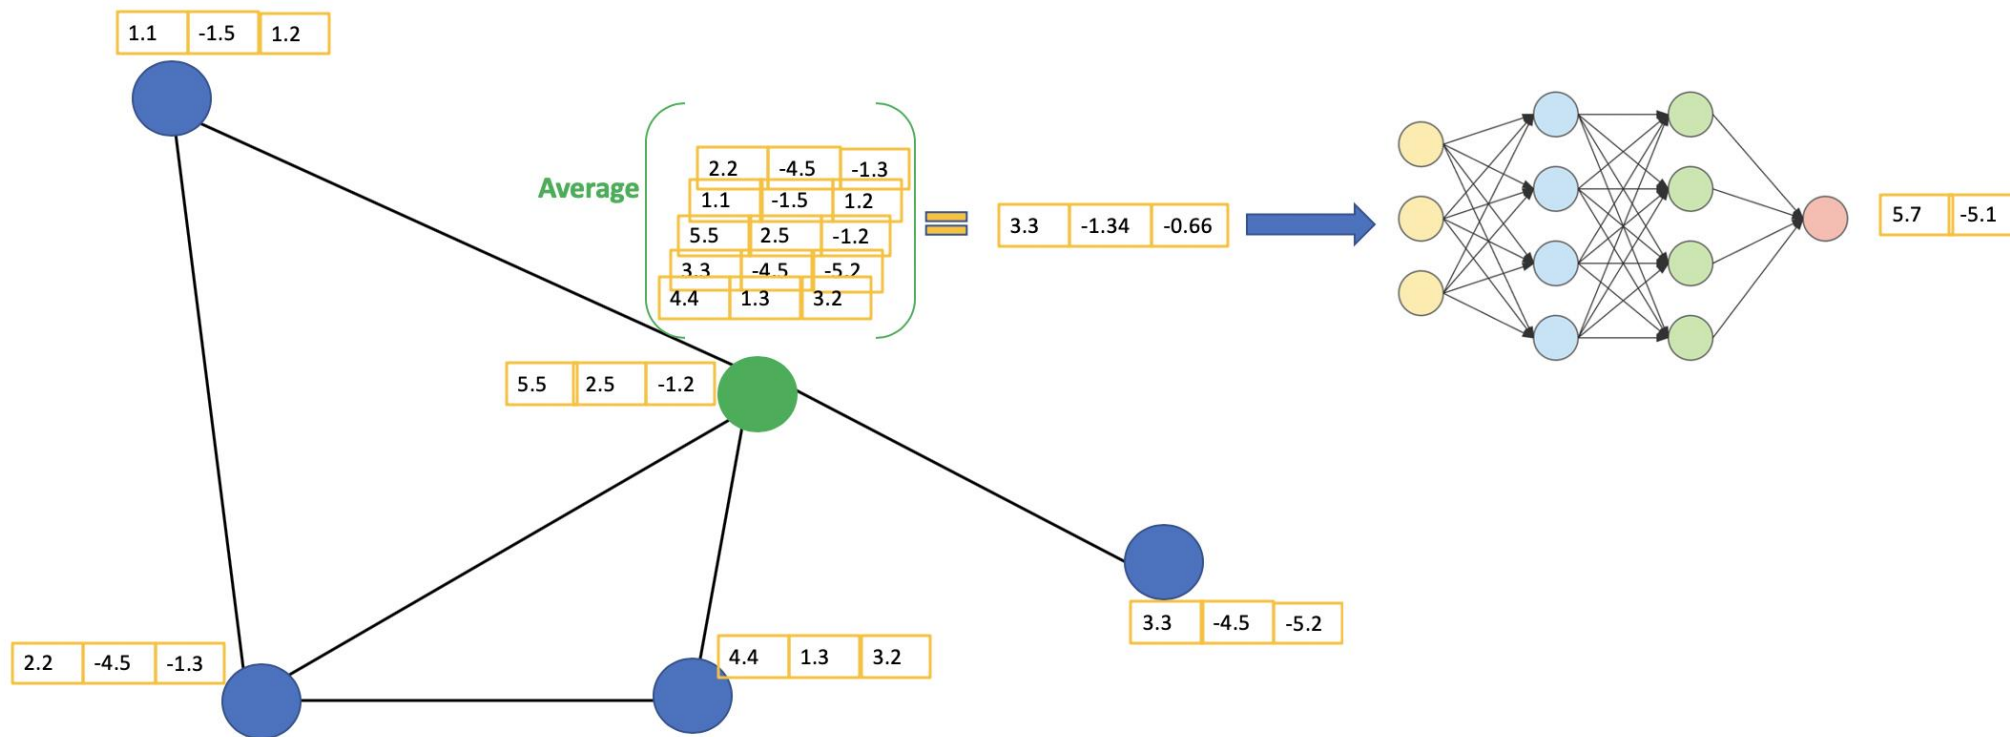


b) 图像分割




d) 引文网络节点分类

Main Idea of GCN




The feature values of the green node's neighbors, including itself, then take the average. The result will be passed through a neural network to return a resulting vector.

Main Idea of GCN



Adjacency List



Adjacency List


```
1  a -> { a b }
2  b -> { c }
3  c -> { d }
4  d -> { b c }
```

Adjacency Matrix

Let G be a graph with " n " nodes that are assumed to be ordered from v_1 to v_n .


The $n \times n$ matrix A is called an **adjacency matrix**, which satisfies:

$$\begin{cases} a_{ij} = 1 & \text{if there exists a path from node } i \text{ to node } j; \\ a_{ij} = 0 & \text{otherwise.} \end{cases}$$



$$\begin{aligned} A_{ij} &= 1 && \text{if there is a link from node } i \text{ to node } j \\ A_{ij} &= 0 && \text{otherwise} \end{aligned}$$

$$A = \begin{pmatrix} 0 & 1 & 0 & 1 \\ 1 & 0 & 0 & 1 \\ 0 & 0 & 0 & 1 \\ 1 & 1 & 1 & 0 \end{pmatrix}$$




$$A = \begin{pmatrix} 0 & 0 & 0 & 1 \\ 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 1 & 1 & 0 \end{pmatrix}$$

Degree Matrix

Let G be a graph with " n " nodes that are assumed to be ordered from v_1 to v_n .

The $n \times n$ matrix D is called a **degree matrix**, which satisfies:

$$\begin{cases} d_{ij} = \deg(v_i) & \text{if } i = j; \\ d_{ij} = 0 & \text{otherwise.} \end{cases}$$




$$\begin{pmatrix} 4 & 0 & 0 & 0 & 0 & 0 \\ 0 & 3 & 0 & 0 & 0 & 0 \\ 0 & 0 & 2 & 0 & 0 & 0 \\ 0 & 0 & 0 & 3 & 0 & 0 \\ 0 & 0 & 0 & 0 & 3 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 \end{pmatrix}$$

Degree Matrix

Where the degree $\deg(v_i)$ of a vertex counts the number of times an edge terminates at that vertex.

In an undirected graph, this means that each **loop** increases the degree of a vertex by **two**.


In a directed graph, the term degree may refer either to indegree (the number of incoming edges at each vertex) or outdegree (the number of outgoing edges at each vertex).




$$\begin{pmatrix} 4 & 0 & 0 & 0 & 0 & 0 \\ 0 & 3 & 0 & 0 & 0 & 0 \\ 0 & 0 & 2 & 0 & 0 & 0 \\ 0 & 0 & 0 & 3 & 0 & 0 \\ 0 & 0 & 0 & 0 & 3 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 \end{pmatrix}$$

Message Passing

What is the relationship between A and B?



Message Passing



	A	B	C	D	E
A	0	0	0	0	1
B	0	0	0	1	1
C	0	0	0	1	1
D	0	1	1	0	1
E	1	1	1	1	0

Adjacency matrix A

	A	B	C	D	E
A	1	0	0	0	0
B	0	2	0	0	0
C	0	0	2	0	0
D	0	0	0	3	0
E	0	0	0	0	4

Degree matrix D

A	-1.1	3.2	4.2
B	0.4	5.1	-1.2
C	1.2	1.3	2.1
D	1.4	-1.2	2.5
E	1.4	2.5	4.5

Feature vector X

Message Passing: Update Adjacency Matrix

$$\tilde{A} = A + \lambda I$$

	A	B	C	D	E
A	0	0	0	0	1
B	0	0	0	1	1
C	0	0	0	1	1
D	0	1	1	0	1
E	1	1	1	1	0

Adjacency matrix A



1	0	0	0	0
0	1	0	0	0
0	0	1	0	0
0	0	0	1	0
0	0	0	0	1

Identity matrix I



	A	B	C	D	E
A	1	0	0	0	1
B	0	1	0	1	1
C	0	0	1	1	1
D	0	1	1	1	1
E	1	1	1	1	1

New Adjacency matrix \tilde{A}

Message Passing: Update Degree Matrix

	A	B	C	D	E
A	1	0	0	0	1
B	0	1	0	1	1
C	0	0	1	1	1
D	0	1	1	1	1
E	1	1	1	1	1

New adjacency matrix \tilde{A}



2	0	0	0	0
0	3	0	0	0
0	0	3	0	0
0	0	0	4	0
0	0	0	0	5

New degree matrix \tilde{D}



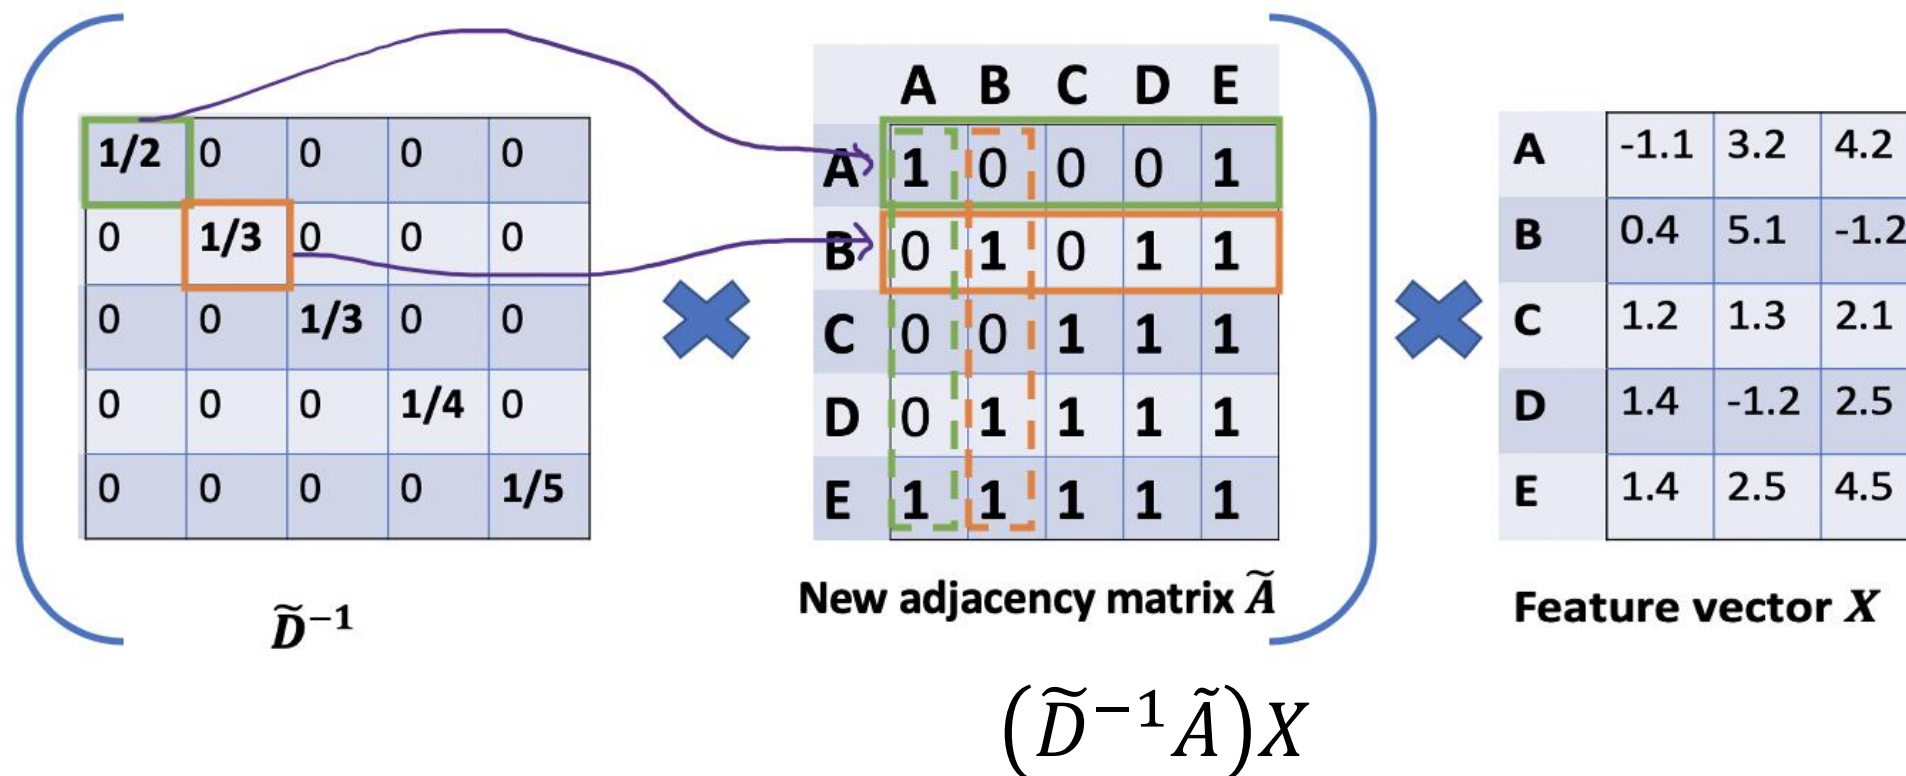
1/2	0	0	0	0
0	1/3	0	0	0
0	0	1/3	0	0
0	0	0	1/4	0
0	0	0	0	1/5

\tilde{D}^{-1}

Message Passing: Scale Factor

We consider $(\tilde{D}^{-1}\tilde{A})X$, so \tilde{D}^{-1} be the **scale factor** of \tilde{A} .

From this perspective, each row i of \tilde{A} will be scaled by \tilde{D}_{ii} .




Message Passing: Scale Factor

From this perspective, each row i of \tilde{A} will be **scaled** by \tilde{D}_{ii} .

\tilde{A} is a **symmetric** matrix, it means row i is the same value as column i .

If we scale each row i of \tilde{A} by \tilde{D}_{ii} , intuitively, we have a feeling that we should do the *same* for its corresponding **column**.




$$\tilde{D}^{-1}\tilde{A}\tilde{D}^{-1}X$$

Message Passing: Scale Factor

The new scaler gives us the “**weighted**” average.

The idea of this weighted average is that we assume **low-degree** nodes would have **bigger impacts** on their neighbors, whereas **high-degree** nodes generate **lower impacts** as they scatter their influence at too many neighbors.



$$\tilde{D}^{-1} \tilde{A} \tilde{D}^{-1} X$$

Message Passing: Scale Factor

When using two scalers (\tilde{D}_{ii} and \tilde{D}_{jj}), we normalize twice: one for row, the other for column.

It will make sense if we rebalance by modifying $\tilde{D}_{ii}\tilde{D}_{jj}$ to $(\tilde{D}_{ii}\tilde{D}_{jj})^{1/2}$

2	0	0	0	0
0	3	0	0	0
0	0	3	0	0
0	0	0	4	0
0	0	0	0	5

\tilde{D}

→

1/2	0	0	0	0
0	1/3	0	0	0
0	0	1/3	0	0
0	0	0	1/4	0
0	0	0	0	1/5

\tilde{D}^{-1}

1/ √2	0	0	0	0
0	1/ √3	0	0	0
0	0	1/ √3	0	0
0	0	0	1/2	0
0	0	0	0	1/ √5

$\tilde{D}^{-1/2}$

$$\tilde{D}^{-\frac{1}{2}} \tilde{A} \tilde{D}^{\frac{1}{2}} X$$

Message Passing: Update Degree Matrix

Quick summary so far:

$\tilde{A}X$: sum of all neighbors' feature vectors, including itself.

$\tilde{D}^{-1} \tilde{A}X$: average of all neighbors' feature vectors (including itself). Adjacency matrix is scaled by rows

$\tilde{D}^{-1/2} \tilde{A} \tilde{D}^{-1/2} X$: average of all neighbors' feature vectors (including itself). The adjacency matrix is scaled by both rows and columns. By doing this, we get the weighted average preferring on low-degree nodes.

Message Passing: Update Degree Matrix

Ok, now let's put things together.

Let's call $\hat{A} = \tilde{D}^{-1/2} \tilde{A} \tilde{D}^{-1/2}$ just for a clear view. With 2-layer GCN, we have the form of our forward model as below.

$$Z = f(X, A) = \text{softmax}\left(\hat{A} \text{ReLU}\left(\hat{A} X W^{(0)}\right) W^{(1)}\right)$$

The diagram illustrates the forward model of a 2-layer Graph Convolutional Network (GCN). The equation is $Z = f(X, A) = \text{softmax}\left(\hat{A} \text{ReLU}\left(\hat{A} X W^{(0)}\right) W^{(1)}\right)$. Annotations include: a purple arrow pointing from the \hat{A} term to the text 'Scaled adjacency matrix (N x N)'; a purple arrow pointing from the X term to the text 'Feature vector matrix (N x C)'; a purple arrow pointing from the $W^{(0)}$ term to the text 'Trainable weights (C x H)'; an orange arrow pointing from the $W^{(1)}$ term to the text 'Trainable weights (H x F)'; and a large purple arrow pointing from the entire expression inside the softmax function to the text 'First layer'.

Recall that N is #nodes, C is #dimensions of feature vectors. We also have H is #nodes in the hidden layer, and F is the dimensions of resulting vectors.


Loss Function

The Loss function is simply calculated by the **cross-entropy error** over all labeled examples.

Y_l is the set of node indices that have labels.

$$\mathcal{L} = - \sum_{l \in \mathcal{Y}_L} \sum_{f=1}^F Y_{lf} \ln Z_{lf}$$




Layers



The number of layers is the farthest distance that node features can travel.

Each node can only get the information from its neighbors with 1-layer GCN.

The gathering information process takes place independently, at the same time for all the nodes.



/02 Benchmarking Graph Neural Networks

GNN的基准化

Benchmarking Graph Neural Networks

Issues:

GNN has been increasingly difficult to gauge the effectiveness of new GNNs and compare models in the **absence** of a standardized benchmark with consistent experimental settings and large datasets.

Solution:

The authors propose a reproducible **GNN benchmarking framework**, with the facility for researchers to add new datasets and models conveniently.

Applications:

The benchmarking framework can be applied to novel medium-scale graph datasets from mathematical modeling, computer vision, chemistry and combinatorial problems to establish key operations when designing effective GNNs.

More:

Precisely, graph convolutions, anisotropic diffusion, residual connections and normalization layers are universal building blocks for developing robust and scalable GNNs.

Benchmarking Graph Neural Networks

Questions Proposed:

- ✓ How to build powerful GNNs has become central?
- ✓ How to study and quantify the impact of theoretical developments for GNNs?

Benchmarking Graph Neural Networks

Contributions::

- I. We release an **open benchmark infrastructure** for GNNs, hosted on GitHub based on PyTorch and DGL libraries;
- II. We aim to go beyond the popular but small CORA and TU datasets by **introducing medium-scale datasets** *with 12k-70k graphs of variable sizes 9-500 nodes*. Proposed datasets are from mathematical modeling (Stochastic Block Models), computer vision (super-pixels), combinatorial optimization (Traveling Salesman Problem) and chemistry (molecules' solubility).
- III. We identify important **building blocks** of GNNs with the proposed benchmark infrastructure. Graph convolutions, anisotropic diffusion, residual connections, and normalization layers stick out as most useful to design efficient GNNs.

Benchmarking Graph Neural Networks


Contributions::

- IV. We fix a parameter budget for all models and analyze performance trends to identify the important **GNN mechanisms**.
- V. The numerical results are entirely **reproducible**. We make it simple to reproduce the reported results by running scripts. Besides, the installation and execution of the benchmark infrastructure are explained in detail in the [GitHub repository](https://github.com/graphdeeplearning/benchmarking-gnns).

Superpixels


A superpixel can be defined as a **group** of pixels that share common characteristics (like pixel intensity).

- Superpixels carry more information than pixels.
- Superpixels have a perceptual meaning since pixels belonging to a given superpixel share similar visual properties.
- Superpixels provide a convenient and compact representation of images that can be very useful for computationally demanding problems.



<https://zhouyifan.net/2020/01/30/DIP-image-segmentation-project/>


Superpixels



Sample images and their superpixel graphs. The graphs of SLIC superpixels (at most 75 nodes for MNIST and 150 nodes for CIFAR10) are 8-nearest neighbor graphs in the Euclidean space and node colors denote the mean pixel intensities.



微信扫一扫关注我们



七月在线
JULYEDU.COM

THANKS

Speaker name and title

<https://www.julyedu.com>